

● OPEN TO WORK · HYBRID / PUNE, INDIA

SENIOR AI ENGINEER · PRODUCTION GENAI, LLMS & AGENTS

Satyabhan Umed Bhadoriya

I build things that don't break — assembly lines, Linux fleets, and now AI systems.

Senior AI Engineer with 8+ YOE (Enterprise & AI). Creator of **CLINIRISE**, a production-ready healthcare WhatsApp AI CRM System powered by Google Gemini, Qdrant, Redis, and GCP Cloud Run. Specialist in Multi-Agent System design, RAG architecture, and AI/GenAI model optimization. Experienced with LangChain, Vertex AI, Groq, AI SDLC tools, and enterprise infrastructure monitoring, as well as authoring over 400 chapters of CI-validated open-source engineering curricula. Adept at applying TypeSafe Jev typed decision models for guardrail enforcement in agentic workflows.

8+

Years — automation → Linux
→ AI/LLM

400+

Chapters authored, 4 CI-
validated curricula

860+

Automated tests on
CLINIRISE

15+

Automation projects
delivered

✉ Email me

Résumé (PDF) ↓

GitHub ↗

LinkedIn ↗

+91 79727 85226

Featured project ↓

Open-source curricula ↓

01

Experience / 2018 — present

○ Independent AI Engineer (Self-Directed)

Pune, India

MAY 2026 — PRESENT

- Lead active development and pilot management of **CLINIRISE** for 2 live clinic clients, continuously iterating on core RAG quality and prompt reliability.
- Architected and deployed comprehensive open-source engineering curricula, building full-scale labs for **eBPF observability**, **vLLM model serving**, and **AI threat defense**.

○ Technical Support Engineer

Aman Security India Pvt. Ltd. — deployed on-site at CBSE Regional Office, Pune

MARCH 2024 — MAY 2026

- **AI Product Development:** Architected and independently built **CLINIRISE** (outside of core support duties) applying GenAI, Python, RAG, and Vertex AI.
- **AI SDLC & Voice Prototyping:** Utilized Claude Code and Antigravity IDE for an AI-assisted development workflow, setting up CI/CD pipelines via GitHub Actions. Prototyped STT/TTS pipelines for future voice-enabled patient interactions.
- **Infrastructure & Monitoring:** Managed virtualized VMware environments, monitoring infrastructure health and service uptime using Prometheus and Grafana dashboards.
- **Technical Operations:** Provided Tier-level technical support and issue resolution using ServiceNow, Jira, and remote access diagnostic tools.

○ Linux Administrator

Oceanic EOPS LLP — remote support for client sites, Pune

2023 — 2024

- Managed RHEL 7 production servers, ensuring high availability and secure infrastructure.
- Applied DISA STIG hardening, executed OpenSCAP compliance auditing, and configured strict SELinux policies.
- Maintained continuous system visibility and monitored performance metrics via Grafana.

○ Application Engineer

Precision Automation, Pune

2018 — 2023

- Delivered end-to-end automation engineering across ~15 major industrial projects, including assembly lines, packaging, and material handling systems.

- Programmed advanced HMI/PLC logic (GX Works / Mitsubishi), consistently maintaining 99.8% operational uptime across deployed systems.


02

Featured project / shipped solo, end to end

CLINIRISE — Autonomous Healthcare CRM & RAG System

Multi-clinic WhatsApp CRM, solo-built end to end — [live app ↗](#)

- Pilot testing — 2 clinic clients



Solo-built, end-to-end multi-clinic CRM currently in pilot testing with 2 live clients. (Live App: [gen-lang-client-0155046437.web.app](#))

- **Multi-Agent System Design:** Architected a supervisor-worker multi-agent topology using LangGraph-style orchestration, including intent-router, scheduling, FAQ RAG, and escalation agents. Implemented typed tool calls, shared agent state, and confidence-gated handoffs for clean agent-to-agent delegation.
- **RAG Architecture & Implementation:** Engineered hybrid retrieval via Qdrant vector search and two-tier Redis caching to achieve sub-2 second response latency. Built semantic intent routing to separate deterministic transactional actions from generative queries, reducing LLM token consumption by ~40%.
- **AI Model Optimization:** Developed multi-tier model resiliency (Gemini Flash → Flash-Lite → template fallback) integrating INT8 quantization, KV-cache optimization, and continuous batching. Benchmarked LLM/STT/TTS stacks against TTFT, error drift, and token cost using structured LLM evaluations.

- **AgentOps & Responsible AI:** Implemented clinical guardrails including PII redaction (DPDP Act compliance), consent management, and human handoff escalation. Applied TypeSafe Jev (Choice/Score/Boolean) for fast, low-cost verification in place of full LLM calls.
- **Cloud & DevOps:** Built with React 18, TypeScript, and Vite frontend alongside a Node.js/Express backend on GCP Cloud Run. Managed infrastructure via Docker, GitHub Actions CI/CD, and Workload Identity Federation. Validated with 860+ automated tests (Vitest, Cypress E2E, k6) targeting a scale of 1,000 clinics processing 100 patients/day.

WhatsApp Business API

Google Gemini

Qdrant RAG

GCP Cloud Run

Express / Node.js

React 18 + TypeScript

Firestore

Redis / Upstash

Firebase Hosting

BigQuery + Looker

Vitest · Cypress · k6

GitHub Actions

03 Skills / what I actually use, daily

AI & Systems Engineering

LLM Deployment, GenAI, RAG Architecture, Tool/Function Calling, Prompt Engineering, and Multi-Agent Systems

Frameworks & Libraries

PyTorch, Hugging Face Transformers, FastAPI, LangChain, LangGraph, and TypeSafe Jev

AI Model Optimization

LoRA/DPO Fine-Tuning, INT8 Quantization, vLLM Serving, and KV-Cache Optimization

AI Platforms & Inference

GCP Vertex AI (Gemini), Groq (Fast Inference), Ollama, Azure STT, and ElevenLabs

AI SDLC Tools

Claude Code, GitHub Copilot, and Antigravity IDE

Vector & Data

Qdrant (Vector DB), Redis, PostgreSQL, and Firebase/Firestore

Cloud, DevOps & Infra

Linux/RHEL 7, Docker, Terraform, Ansible, Prometheus, Grafana, GitHub Actions, VMware, OpenSCAP, and STIG Hardening

Languages

Python (Primary), TypeScript, and Node.js

04 Open-source engineering / 400+ chapters, CI-validated

Author & maintainer of three production-depth engineering curricula — each its own live site with runnable labs and automated CI. Built outside client work, in the open, continuously growing.

AI / ML ENGINEERING

● 9 parts · 138 ch

AI-ML

From-scratch autograd, BPE tokenizers, scaled dot-product attention and embeddings — then production system design, deployment, observability and enterprise-scale AI. Parts 1–5 build the primitives in pure Python; Parts 6–9 operate the standard production stack.

Explore the curriculum →

69/69 automated CI checks
· zero dependencies · 6 live browser labs

CLOUD & DEVOPS

● 5 parts · 88 ch

DevOps-Infra

Infrastructure as Code, Terraform, Docker and Kubernetes from first principles — then IaC at scale, platform engineering, cloud networking & identity, and observability & incident engineering. Runnable labs, 3 a.m. production scenarios and SRE playbooks.

Explore the curriculum →

4-gate CI (validate + terraform + docker + kubeconform) · 14 deep-dive guides · 20+ lab files

LINUX, HPC & SECURITY

● 21 modules · 126 lessons

linux-hpc-security

Linux internals and DISA STIG hardening, HPC benchmarking, eBPF, Ansible, Warewulf, Slurm, RDMA/InfiniBand fabric, parallel storage, kernel tuning, SOC/NOC operations, Splunk SIEM and red/blue/purple teaming — modelled in runnable pure-Python labs.

Explore the curriculum →

63 runnable labs · 800+ interview Q&A · 5 course parts

🔒 AI THREAT DEFENSE

● Production Grade

AI Threat Defense

Production-grade AI
Threat Detection,
Adversarial Robustness &
Go Gateway Engineering
Curriculum.

Explore the
curriculum →

Live Curriculum

★ REPOSITORIES & VERIFIED CI

- **AI-ML & AI-Lab:** Authored a 9-course curriculum covering scalar autograd, BPE tokenizers, attention mechanisms, RAG + LangGraph/ReAct agent loops, LoRA/DPO fine-tuning, INT8 quantization, KV-cache, and vLLM serving. Maintains 69/69 passing CI checks.
- **AI Threat Defense:** Developed production-grade training materials for AI Threat Detection, Adversarial Robustness, and Go Gateway Engineering.
- **DevOps-Infra:** Built extensive SRE playbooks and infrastructure code utilizing Terraform IaC and Docker multi-stage hardened builds, covering SLI/SLOs, golden signals, and incident runbooks secured by a 4-gate CI pipeline.
- **Linux-HPC-Security:** Created comprehensive systems engineering modules from the metal up, detailing Linux internals, hardening, DISA STIGs, Ansible, Slurm HPC, and featuring an advanced eBPF Observability Capstone.

Always building. Always shipping. Always learning in public.

Profile

Senior AI Engineer with 8+ YOE (Enterprise & AI). Creator of **CLINIRISE**, a production-ready healthcare WhatsApp AI CRM

System powered by Google Gemini, Qdrant, Redis, and GCP Cloud Run. Specialist in Multi-Agent System design, RAG architecture, and AI/GenAI model optimization. Experienced with LangChain, Vertex AI, Groq, AI SDLC tools, and enterprise infrastructure monitoring, as well as authoring over 400 chapters of CI-validated open-source engineering curricula. Adept at applying TypeScript Jev typed decision models for guardrail enforcement in agentic workflows.

Education

B.E. Mechanical Engineering — First Class
Savitribai Phule Pune University, Dhole Patil College of Engineering (DPCOE), 2014

[Email](#) [Call](#) [Résumé \(PDF\)](#) [GitHub](#) [LinkedIn](#) [Back to top ↑](#)

VISITOR METRICS

VISITS 0

YOUR IP ADDRESS

223.236.99.165