

SATYABHAN UMED BHADORIYA

AI Engineer | Multi-Agent Systems | RAG Architecture | LangChain/LangGraph | Production GenAI

Pune, India | +91 79727 85226 | satyabhan.bhadoriya@gmail.com

satyabhan007.github.io | [GitHub](#) | [LinkedIn](#) | Immediate Joiner • Hybrid/Pune

PROFESSIONAL SUMMARY

AI Engineer with 2+ years building and deploying production **LLMs**, **GenAI**, and **Agentic AI** solutions, plus 8+ years of enterprise engineering discipline. Specialist in **Multi-Agent System design**, **RAG architecture & implementation**, and **AI/GenAI model optimization**. Creator of **CLINIRISE** — a production multi-agent WhatsApp CRM powered by Google Gemini, Qdrant, Redis, and GCP Cloud Run — solo-built end-to-end with **860+ automated tests**, in active pilot with 2 clinic clients. Applies **TypeSafe Jev** typed decision models for guardrail enforcement in agentic workflows. Experienced with **LangChain**, **Gemini / Vertex AI**, **Groq**, **AI SDLC tools** (Claude Code, GitHub Copilot), and enterprise infrastructure monitoring. Author of **400+ chapters** of CI-validated open-source engineering curricula.

FEATURED PROJECT — CLINIRISE (PILOT TESTING, 2 CLINIC CLIENTS)

Production Multi-Agent Healthcare CRM & RAG System — solo-built, end-to-end

Live app → gen-lang-client-0155046437.web.app

- **Multi-Agent System Design:** Architected a supervisor-worker multi-agent topology using **LangGraph**-style orchestration — intent-router agent, scheduling agent, FAQ RAG agent, and escalation agent — with typed tool calls, shared agent state, and confidence-gated handoffs for clean agent-to-agent delegation.
- **RAG Architecture & Implementation:** Hybrid retrieval via **Qdrant** vector search + two-tier **Redis** caching, targeting sub-2-second response latency. Semantic intent routing separates deterministic transactional actions from generative queries, reducing LLM token consumption by ~40%.
- **AI Model Optimization:** Multi-tier model resiliency (Gemini Flash → Flash-Lite → template fallback) with **INT8 quantization**, **KV-cache** optimization, and **continuous batching**. Benchmarked LLM/STT/TTS stacks against TTFT, error drift, and token cost using structured **LLM Evals**.
- **AgentOps & Responsible AI:** Clinical guardrails with PII redaction (DPDP Act), consent management, and human handoff escalation. Applied **TypeSafe Jev** (Choice/Score/Boolean) for fast, low-cost verification in place of full LLM calls.
- **Cloud & DevOps:** React 18 + TypeScript + Vite frontend, Node.js/Express on **GCP Cloud Run**, **Docker**, **GitHub Actions** CI/CD, Workload Identity Federation. **860+ automated tests** (Vitest + Cypress E2E + k6). Scale target: 1,000 clinics × 100 patients/day.

EXPERIENCE

Technical Support Engineer | Aman Security India Pvt. Ltd.

March 2024 – May 2026 | Pune, India

- **Self-Initiated AI Product Development:** Architected and independently built **CLINIRISE**, a production **LLM-powered multi-clinic WhatsApp CRM** (outside core support duties) applying **Gen AI**, **Python**, **RAG**, and **Vertex AI (Gemini)**. Currently in pilot testing with real clinic clients.
- **AI SDLC & Voice Prototyping:** Used **Claude Code** and **Antigravity IDE** as an AI-assisted development workflow to build and debug the platform, utilizing **GitHub Actions** for CI/CD. Prototyped and explored **STT/TTS** pipelines for future voice-enabled patient interactions.
- **Infrastructure & Monitoring:** Managed virtualized environments (VMware) and monitored infrastructure health and service uptime using **Prometheus** and **Grafana** dashboards.
- **Technical Operations:** Provided technical support and issue resolution using ServiceNow and Jira, diagnosing and fixing system issues via remote access tools.

Independent AI Engineer (Self-Directed)

May 2026 – Present | Pune, India

- Continued active development and pilot management of **CLINIRISE** (2 live clinic clients); iterating on core RAG quality and prompt reliability.
- Architected and deployed comprehensive open-source engineering curricula, building full-scale labs for **eBPF observability**, **vLLM model serving**, and **AI threat defense**.

Linux Administrator | Oceanic EOPS LLP

2023 – 2024 | Pune, India

- Managed RHEL 7 production servers; applied **DISA STIG hardening**, OpenSCAP compliance auditing, SELinux policies; monitored with Grafana.

Application Engineer | Precision Automation

2018 – 2023 | Pune, India

- Delivered end-to-end automation engineering across ~15 projects (assembly lines, packaging, material handling). Programmed HMI/PLC logic (GX Works / Mitsubishi). 99.8% operational uptime.

SKILLS

- **AI & Systems Engineering:** LLM Deployment, GenAI, RAG Architecture, Tool/Function Calling, Prompt Engineering, Multi-Agent Systems
- **Frameworks & Libraries:** PyTorch, Hugging Face Transformers, FastAPI, LangChain, TypeSafe Jev
- **AI Model Optimization:** LoRA/DPO Fine-Tuning, INT8 Quantization, vLLM Serving, KV-Cache
- **AI Platforms & Inference:** GCP Vertex AI (Gemini), Groq (Fast Inference), Ollama, Azure STT, ElevenLabs
- **AI SDLC Tools:** Claude Code, GitHub Copilot, Antigravity IDE
- **Vector & Data:** Qdrant (Vector DB), Redis, PostgreSQL, Firebase/Firestore
- **Cloud, DevOps & Infra:** Linux/RHEL 7, Docker, Terraform, Ansible, Prometheus, Grafana, GitHub Actions, VMware, OpenSCAP, STIG Hardening
- **Languages:** Python (primary), TypeScript, Node.js

OPEN-SOURCE ENGINEERING & BENCHMARKING

Author & Maintainer — Engineering Curricula (4 Live Repos, CI-Validated)

satyabhan007.github.io

- **AI-ML & AI-Lab** (satyabhan007.github.io/ai-lab): 9-course curriculum — scalar autograd, BPE tokenizer, attention, **RAG + LangGraph/ReAct agent loops, LoRA/DPO fine-tuning, INT8 quantization, KV-cache, vLLM serving. 69/69 CI checks passing.**
- **AI Threat Defense** (satyabhan007.github.io/ai-threat-defense): Production-grade AI Threat Detection, Adversarial Robustness & Go Gateway Engineering Curriculum.
- **DevOps-Infra** (satyabhan007.github.io/DevOps-Infra): Terraform IaC, **Docker** multi-stage hardened builds, SRE playbooks (SLI/SLO, golden signals, incident runbooks). 4-gate CI pipeline.
- **Linux-HPC-Security** (satyabhan007.github.io/linux-hpc-security): Systems engineering from the metal up — Linux internals, hardening, DISA STIGs, Ansible, Slurm HPC. Features an advanced [eBPF Observability Capstone](#).

EDUCATION

Bachelor of Engineering (B.E.), Mechanical Engineering — First Class | 2014

Savitribai Phule Pune University — Dhole Patil College of Engineering (DPCOE)